

# Kevin Wu

PHD CANDIDATE · STANFORD UNIVERSITY

735 Campus Dr. Stanford, CA 94305

✉ kevinwu@stanford.edu | 🏠 www.kevinwu.ai | 📧 kevin-y-wu

## Education

### Stanford University

PH.D. BIOMEDICAL DATA SCIENCE

• Advisor: James Zou & Daniel E. Ho

Stanford, CA

09/2020 - present

### Harvard University

M.ENG. COMPUTATIONAL SCIENCE AND ENGINEERING

• Advisors: Gabriel Kreiman & David Cox

Cambridge, MA

08/2016 - 05/2018

### Duke University

B.A. STATISTICS

Durham, NC

08/2011 - 05/2015

## Publications

### PREPRINT

1. Kevin Wu\*, Angela Zhang\*, Joshua Guild, Mert Yuksekgonul, Eric Wu, Joseph C. Wu, James Zou. *Elucidating the Mechanisms of Gender and Racial Bias of Large Language Models in Clinical Management of Cardiovascular Disease, Pain, and Cancer Screening and Treatment*. **Cell Reports Medicine (Under Review)**. 2024.
2. Kevin Wu\*, Eric Wu\*, James Zou. *FineTuneBench: How well do commercial fine-tuning APIs infuse knowledge into LLMs?*. **NEJM AI (Revise & Resubmit)** 2025.
3. Andy Zhou, Kevin Wu, Yi Zeng, Yu Yang, Shuang Yang, Sanmi Koyejo, James Zou, Bo Li. *AutoRedTeamer: Automated and Adaptive Red Teaming Agent against Language Models*. 2024.

### PUBLISHED

1. Kevin Wu\*, Eric Wu\*, Ally Cassasola, Angela Zhang, Kevin Wei, Teresa Nguyen, Sith Riantawan, Patricia Shi Riantawan, Daniel E Ho, James Zou. *How well do LLMs cite relevant medical references? An evaluation framework and analyses*. **Nature Communications**, 2025.
2. Lucas Rodrigues\*, Kevin Wu\*, Garrett Harvey, Gerald Post, Aubrey Miller, Abigail Hull, Lindsay Lambert, Christina Lopes, James Zou. *Real-world evidence couples genomic biomarkers with therapeutic outcomes for canine hemangiosarcoma*. **Scientific Reports**, 2025.
3. Kevin Wu\*, Eric Wu\*, and James Zou. *ClashEval: Quantifying the tug-of-war between an LLM's internal prior and external evidence*. **NeurIPS**, 2024.
4. Kevin Wu, Eric Wu, Kit Rodolfa, Daniel E Ho, James Zou. *Regulating AI Adaptation: An Analysis of AI Medical Device Updates*. Conference on Health, Inference, and Learning (**CHIL**), 2024.
5. Kevin Wu, Eric Wu, Brandon Theodorou, Weixin Liang, Christina Mack, Lucas Glass, Jimeng Sun, James Zou. *Characterizing the clinical adoption of medical AI devices through US insurance claims*. **New England Journal of Medicine AI**, 2024.
6. Yongchan Kwon\*, Eric Wu\*, Kevin Wu\*, James Zou. *DataInf: Efficiently estimating data influence in LoRA-tuned LLMs and diffusion models*. **ICLR**, 2023.
7. Kevin Wu, Dominik Dahlem, Christopher Hane, Eran Halperin, James Zou. *Collecting data when missingness is unknown: a method for improving model performance given under-reporting in patient populations*, Conference on Health, Inference, and Learning (**CHIL**), 2023.

8. Kevin Wu, Lucas Rodrigues, Gerald Post, Garrett Harvey, Michelle White, Aubrey Miller, Lindsay Lambert, Benjamin Lewis, Christina Lopes, James Zou. *Analyses of canine cancer mutations and treatment outcomes using real-world clinico-genomics data of 2119 dogs*, **npj Precision Oncology**, 2023.
9. Kevin Wu, Eric Wu, Michael DAndrea, Nandini Chitale, Melody Lim, Marek Dabrowski, Klaudia Kantor, Hanoor Rang, Ruis-han Liu, Marius Garmhausen, Navdeep Pal, Chris Harbron, Shemra Rizzo, Ryan Copping, James Zou. *Machine learning prediction of clinical trial operational efficiency*, **AAPSJ**, 2022.
10. Kevin Wu\*, Eric Wu\*, James Zou. *Explaining medical AI performance disparities across sites with confounder Shapley value analysis*, **ML4H**, 2021.
11. Eric Wu, Kevin Wu, Roxana Daneshjou, David Ouyang, Daniel E Ho, James Zou. *How medical AI devices are evaluated: limitations and recommendations from an analysis of FDA approvals*, **Nature Medicine**, 2021.
12. William Lotter, Abdul Rahman Diab, Bryan Haslam, Jiye G Kim, Giorgia Grisot, Eric Wu, Kevin Wu, Jorge Onieva Onieva, Yun Boyer, Jerrold L Boxerman, Meiyun Wang, Mack Bandler, Gopal R Vijayaraghavan, A Gregory Sorensen, *Robust breast cancer detection in mammography and digital breast tomosynthesis using an annotation-efficient deep learning approach*, **Nature Medicine**, 2021.
13. Kevin Wu, Eric Wu, Yaping Wu, Hongna Tan, Greg Sorensen, Meiyun Wang, Bill Lotter, *Validation of a deep learning mam-mography model in a population with low screening rates*, **NeurIPS Fair ML for Health Workshop**, 2019.
14. Kevin Wu, Eric Wu, Gabriel Kreiman. *Learning scene gist with convolutional neural networks to improve object recognition*, Annual Conference on Information Sciences and Systems (**CISS**), 2018.
15. Eric Wu, Kevin Wu, David Cox, William Lotter, *Conditional infilling GANs for data augmentation in mammogram classifica-tion*, **MICCAI Workshop**, 2018.
16. Brett Walenz, Y Wu, S Song, Emre Sonmez, Eric Wu, Kevin Wu, Pankaj K Agarwal, Jun Yang, Naeemul Hassan, Afroza Sultana, Gensheng Zhang, Chengkai Li, Cong Yu. *Finding, monitoring, and checking claims computationally based on structured data*, **Computation+ Journalism Symposium**, 2014.

## Professional Experience

---

2021 - present	<b>Research Scientist</b> , FidoCure (Part-Time)
2021 - present	<b>Instructor</b> , UpLimit (Part-Time)
2022 - 2023	<b>Graduate Research Intern</b> , Optum Labs (United Healthcare)
2018 - 2020	<b>Machine Learning Engineer</b> , DeepHealth, Acq. by RadNet, Inc in 2020
2017 - 2017	<b>Intern</b> , Waymo
2015 - 2016	<b>Product Manager</b> , Microsoft

## Awards, Fellowships, & Grants

---

2022 - 2024	<b>Stanford Data Science Fellowship</b> , Stanford University	PhD funding
2024	<b>Stanford HAI GCP Grant</b> , Data attribution and design for large language models	\$ 20,000
2023	<b>Stanford HAI GCP Grant</b> , Learning multi-modal, multi-disease representations of medical images	\$ 20,000

## Presentations

---

### INVITED TALKS

Summer 2024. *How medical AI are evaluated, deployed, and updated*. Invited talk: Towards Reliable, Valid, and Safe Systems for Biomedical Data Science, JSM 2024, Portland, OR.

Winter 2023. *Regulating Medical AI*. Invited Talk: AI & Health Regulatory Policy Conference, MIT Jameel Clinic.

### ORAL PRESENTATIONS

*Collecting data when missingness is unknown: a method for improving model performance given under-reporting in patient populations*. CHIL 2023.

*Concordance between dogs and humans: The use of AI in evaluating clinical cancer genomic datasets*. AACR 2021.

## Teaching Experience

---

Winter 2023	<b>BIOMEDIN 202: Biomedical Data Science</b> , Teaching Assistant	<i>Stanford</i>
Spring 2022	<b>CS 235: Biomedical Image Analysis and Interpretation</b> , Teaching Assistant	<i>Stanford</i>
Winter 2018	<b>CS 109b: Introduction to Data Science</b> , Teaching Fellow	<i>Harvard</i>
Winter 2027	<b>CS 207: Systems Development for Computational Science</b> , Teaching Fellow	<i>Harvard</i>
Spring 2015	<b>CS 290L: Everything Data</b> , Undergraduate Teaching Assistant	<i>Duke</i>